

# Thesis: Causal Analysis

*Laurens van der Tas*

## Initialize: Load R environment, data and packages

In this document we will investigate the causal relationships in our dataset using regression analysis.

```
#Set working directory
#setwd("C:/Users/Laurens/Dropbox/University/Year 4/Period 2/Applied Economics Research Course/Thesis/da

#Load required packages
library(foreign)
library(stargazer)
library(ggplot2)
library(aod)
library(gridExtra)
library(ggthemes)
library(dplyr)
library(mfx)
library(corrplot)
library(car)
library(LogisticDx)
library(rms)

#Enable anti-aliasing for rendered graphics
library(knitr)
#opts_chunk$set(out.width = '\\maxwidth')
#dev = "CairoPNG",

#Load dataset
data.dropout <- read.dta("DatasetTrimmed.dta")
  #Read name vector of dataset
  names(data.dropout)

#Factorize binaries
data.dropout$geslachtBin <- factor(data.dropout$geslachtBin, labels = c("Female", "Male"))
data.dropout$allochtoonBin <- factor(data.dropout$allochtoonBin, labels = c("No", "Yes"))
```

## Regression Models

In this section all the investigated models, as well as the obsolete ones, are listed.

### LPM Model

For testing purposes.

```
#TRIVIAL: LPM model of Big Five on dropout
lpm.1 <- lm(vrv_1 ~ zorgvuldig + stabiel + open + extravert + altruisme + allochtoonBin + geslachtBin,
  data = data.dropout)
#Display results
summary(lpm.1)
```

## First Logit Model

Neglects the theory in the sense that it disregards *gender* and *foreign origin* as control variables.

```
logit.1 <- glm(vrv_1 ~ stabiel + open + zorgvuldig + extravert + altrusme + age_opleiding,
               family = binomial(logit),
               data = data.dropout)
```

```
#Display results
summary(logit.1)
```

```
##
## Call:
## glm(formula = vrv_1 ~ stabiel + open + zorgvuldig + extravert +
##      altrusme + age_opleiding, family = binomial(logit), data = data.dropout)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8954  -1.2830   0.8695   0.9972   1.2225
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.02021    0.53107   0.038 0.969644
## stabiel        -0.03668    0.08695  -0.422 0.673153
## open          -0.01039    0.10280  -0.101 0.919523
## zorgvuldig    -0.20678    0.09600  -2.154 0.031244 *
## extravert      0.07932    0.10499   0.755 0.449957
## altrusme       0.01165    0.12152   0.096 0.923644
## age_opleiding  0.05571    0.01645   3.386 0.000708 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 659.08  on 498  degrees of freedom
## Residual deviance: 639.80  on 492  degrees of freedom
## AIC: 653.8
##
## Number of Fisher Scoring iterations: 4
```

```
#Calculate odds ratio vector
coef.vector.1 <- exp(logit.1$coef)

#R squared and Wald test significance
require(rms)
lrm(logit.1)
```

```
## Logistic Regression Model
```

```
##
## lrm(formula = logit.1)
##
##              Model Likelihood      Discrimination      Rank Discrim.
##              Ratio Test              Indexes              Indexes
## Obs          499    LR chi2      19.28    R2          0.052    C          0.599
## 0            186    d.f.           6      g          0.467    Dxy         0.199
## 1            313    Pr(> chi2) 0.0037    gr          1.595    gamma        0.200
```

```
## max |deriv| 7e-06                                gp      0.098    tau-a    0.093
##                                                    Brier    0.226
##
##           Coef      S.E.    Wald Z Pr(>|Z|)
## Intercept      0.0202 0.5311  0.04 0.9696
## stabiel        -0.0367 0.0870 -0.42 0.6732
## open           -0.0104 0.1028 -0.10 0.9195
## zorgvuldig     -0.2068 0.0960 -2.15 0.0312
## extravert       0.0793 0.1050  0.76 0.4500
## altruisme       0.0116 0.1215  0.10 0.9236
## age_opleiding  0.0557 0.0165  3.39 0.0007
##
#Marginal effects
require(mfx)
mfx.1 <- logitmfx(logit.1, data = data.dropout)

#Variance inflation factor: multicollinearity test
require(car)
vif(logit.1)

##           stabiel           open    zorgvuldig    extravert    altruisme
##           1.183144        1.257312        1.223593        1.428105        1.182988
## age_opleiding
##           1.036169

sqrt(vif(logit.1)) > 2 #Bigger than 2 (or sometimes 2.5) signals relatively high multicollinearity


##           stabiel           open    zorgvuldig    extravert    altruisme
##           FALSE          FALSE          FALSE          FALSE          FALSE
## age_opleiding
##           FALSE

#Test for outliers
outlierTest(logit.1)

##
## No Studentized residuals with Bonferonni p < 0.05
## Largest |rstudent|:
##           rstudent unadjusted p-value Bonferonni p
## 406 -1.929694          0.053645          NA

#Cook's distance
#plot(logit.1, which = 4, main = "Model One")

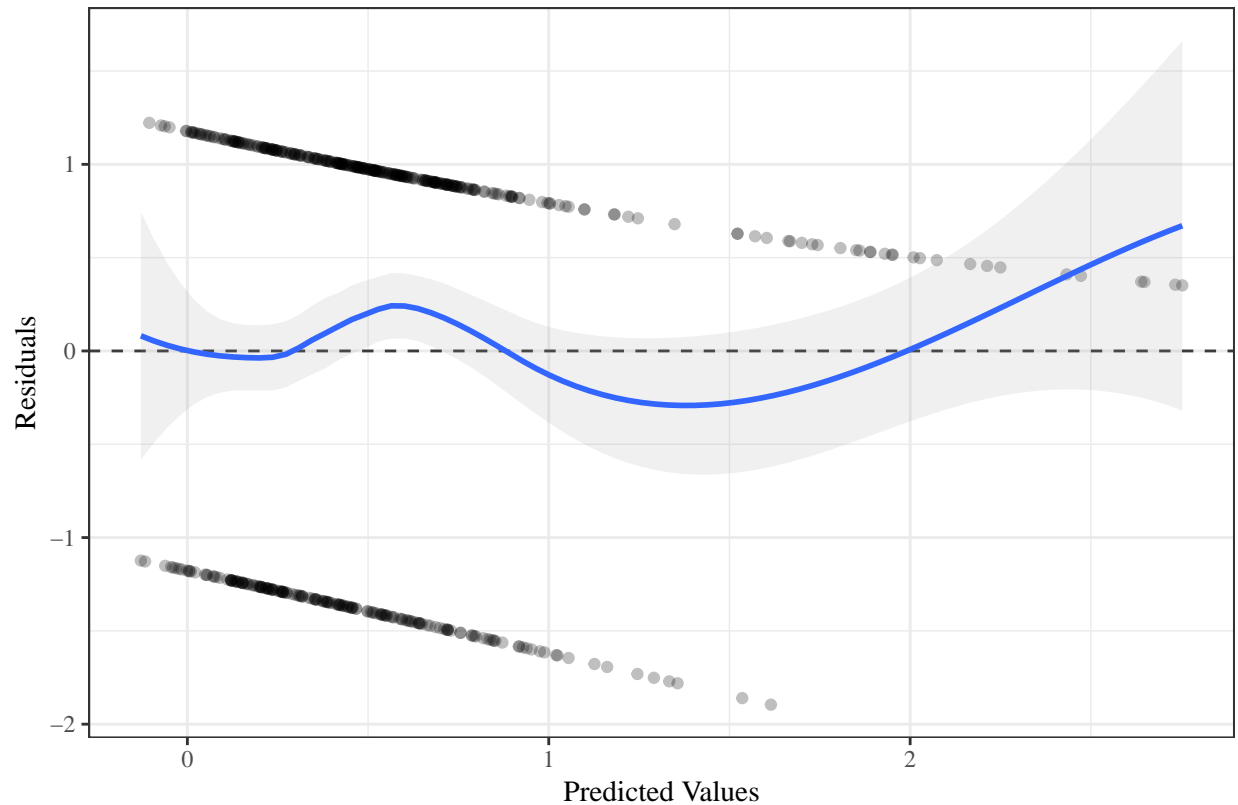
ggplot(aes(x = seq_along(.cooks), y = .cooks), data = logit.1) +
  geom_bar(stat = "identity") +
  theme_bw(base_family = "serif") +
  labs(title = "Model One: Cook's Distance",
       x = "Observation Number",
       y = "Cook's Distance") +
  geom_hline(yintercept = 0.0081, colour = 2, alpha = 0.75, linetype = 5)
```



```
#Resid vs fitted
#plot(logit.1, which = 1, main = "Model One")

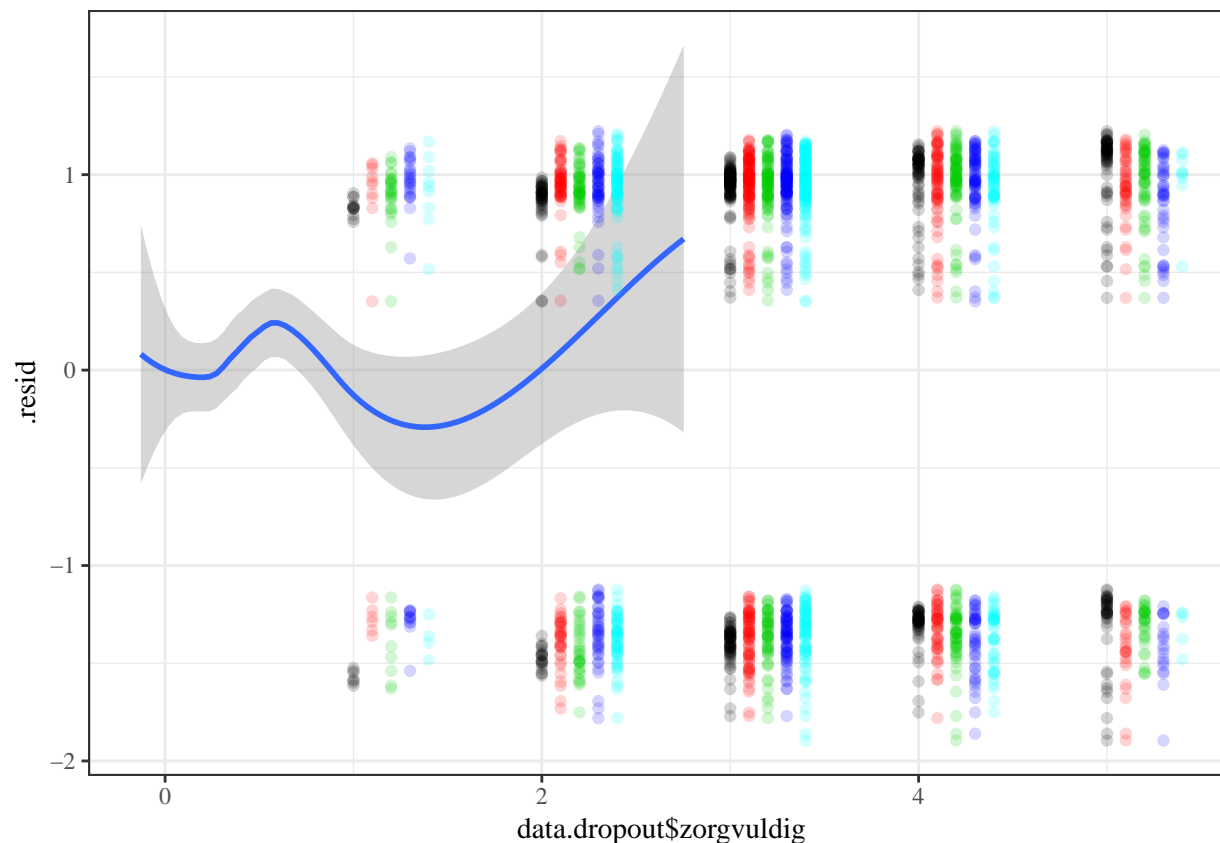
ggplot(aes(x = .fitted, y = .resid), data = logit.1) +
  geom_point(alpha = 1/4) +
  geom_hline(yintercept = 0, linetype = 2, alpha = 0.75) +
  geom_smooth(alpha = 0.15, method = "loess") +
  theme_bw(base_family = "serif") +
  labs(title = "Model One: Residuals vs. Fitted",
       x = "Predicted Values",
       y = "Residuals")
```

Model One: Residuals vs. Fitted



```
#Model fitted on covariates
ggplot(aes(x = data.dropout$zorgvuldig, y = .resid), data = logit.1) +
  theme_bw(base_family = "serif") +
  geom_smooth(aes(x = .fitted, y = .resid)) +
  geom_point(alpha = 1/6) +
  geom_point(aes(x = data.dropout$open+0.1, y = .resid), colour = 2, alpha = 1/6) +
  geom_point(aes(x = data.dropout$stabiel+0.2, y = .resid), colour = 3, alpha = 1/6) +
  geom_point(aes(x = data.dropout$extravert+0.3, y = .resid), colour = 4, alpha = 1/6) +
  geom_point(aes(x = data.dropout$altruisme+0.4, y = .resid), colour = 5, alpha = 1/6)

## `geom_smooth()` using method = 'loess'
```



## Second Logit Model

Added *gender* and *foreign origin* as control variables to the first logit model.

```
logit.2 <- glm(vrv_1 ~ stabiel + open + zorgvuldig + extravert + altruisme + age_opleiding + geslachtBin
               allochtoonBin,
               family = binomial(logit),
               data = data.dropout)

#Display results
summary(logit.2)
```

```
##
## Call:
## glm(formula = vrv_1 ~ stabiel + open + zorgvuldig + extravert +
##       altruisme + age_opleiding + geslachtBin + allochtoonBin, family = binomial(logit),
##       data = data.dropout)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8665  -1.2884   0.8566   0.9997   1.2473
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -0.04660    0.53913  -0.086  0.931121
## stabiel      -0.05222    0.08969  -0.582  0.560368
```

```

## open          -0.02572    0.10521   -0.244  0.806885
## zorgvuldig    -0.19280    0.09915   -1.944  0.051836 .
## extravert      0.08508    0.10532    0.808  0.419199
## altruisme      0.01471    0.12168    0.121  0.903769
## age_opleiding  0.05646    0.01652    3.418  0.000631 ***
## geslachtBinMale 0.15153    0.20361    0.744  0.456721
## allochtoonBinYes 0.06659    0.22302    0.299  0.765271
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 659.08 on 498 degrees of freedom
## Residual deviance: 639.19 on 490 degrees of freedom
## AIC: 657.19
##
## Number of Fisher Scoring iterations: 4
#R squared and Wald test significance
require(rms)
lrn(logit.2)

## Logistic Regression Model
##
## lrn(formula = logit.2)
##
##              Model Likelihood   Discrimination   Rank Discrim.
##              Ratio Test         Indexes         Indexes
## Obs          499   LR chi2      19.89   R2          0.053   C          0.601
## 0             186   d.f.         8      g           0.476   Dxy         0.201
## 1             313   Pr(> chi2) 0.0108   gr           1.610   gamma      0.203
## max |deriv| 8e-06                gp           0.100   tau-a      0.094
##              Brier           0.226
##
##              Coef    S.E.   Wald Z Pr(>|Z|)
## Intercept      -0.0466 0.5391 -0.09 0.9311
## stabiel        -0.0522 0.0897 -0.58 0.5604
## open           -0.0257 0.1052 -0.24 0.8069
## zorgvuldig     -0.1928 0.0992 -1.94 0.0518
## extravert       0.0851 0.1053  0.81 0.4192
## altruisme       0.0147 0.1217  0.12 0.9038
## age_opleiding   0.0565 0.0165  3.42 0.0006
## geslachtBin=Male 0.1515 0.2036  0.74 0.4567
## allochtoonBin=Yes 0.0666 0.2230  0.30 0.7653
##
#Marginal effects
require(mfx)
logitmfx(logit.2, data = data.dropout)

## Call:
## logitmfx(formula = logit.2, data = data.dropout)
##
## Marginal Effects:
##              dF/dx   Std. Err.      z    P>|z|
## stabiel        -0.0120952  0.0207698 -0.5823 0.5603351

```

```
## open -0.0059565 0.0243658 -0.2445 0.8068734
## zorgvuldig -0.0446532 0.0229390 -1.9466 0.0515822 .
## extravert 0.0197039 0.0243864 0.8080 0.4190991
## altruisme 0.0034071 0.0281810 0.1209 0.9037686
## age_opleiding 0.0130772 0.0037752 3.4640 0.0005322 ***
## geslachtBinMale 0.0349592 0.0467581 0.7477 0.4546648
## allochtoonBinYes 0.0153493 0.0511618 0.3000 0.7641653
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## dF/dx is for discrete change for the following variables:
##
## [1] "geslachtBinMale" "allochtoonBinYes"
```

```
#Variance inflation factor: multicollinearity test
require(car)
vif(logit.2)
```

```
##          stabiël          open          zorgvuldig          extravert
##          1.257415          1.315879          1.303901          1.436032
##          altruisme    age_opleiding    geslachtBinMale    allochtoonBinYes
##          1.183412          1.040891          1.124696          1.054038
```

```
sqrt(vif(logit.2)) > 2
```

```
##          stabiël          open          zorgvuldig          extravert
##          FALSE          FALSE          FALSE          FALSE
##          altruisme    age_opleiding    geslachtBinMale    allochtoonBinYes
##          FALSE          FALSE          FALSE          FALSE
```

```
#Test for outliers
outlierTest(logit.2)
```

```
##
## No Studentized residuals with Bonferonni p < 0.05
## Largest |rstudent|:
##          rstudent unadjusted p-value Bonferonni p
## 406 -1.902455          0.057112          NA
```

### Third Logit Model

Added interaction effect between emotional *stability* and *conscientiousness*, as suggested by the psychologist referred to by Yolanda.

```
logit.3 <- glm(vrv_1 ~ stabiël + open + zorgvuldig + extravert + altruisme + age_opleiding + stabiël:zorgvuldig,
               family = binomial(logit),
               data = data.dropout)
```

```
#Display results
summary(logit.3)
```

```
##
## Call:
## glm(formula = vrv_1 ~ stabiël + open + zorgvuldig + extravert +
##       altruisme + age_opleiding + stabiël:zorgvuldig, family = binomial(logit),
##       data = data.dropout)
##
```



```
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9140  -1.2881   0.8516   1.0022   1.2020
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.859716   0.989461   0.869 0.384917
## stabel        -0.290425   0.265614  -1.093 0.274213
## open          -0.020262   0.103352  -0.196 0.844575
## zorgvuldig    -0.441306   0.251862  -1.752 0.079744 .
## extravert      0.076597   0.105109   0.729 0.466160
## altruisme      0.006072   0.121729   0.050 0.960217
## age_opleiding  0.055556   0.016508   3.365 0.000764 ***
## stabel:zorgvuldig 0.073843   0.072905   1.013 0.311126
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 659.08  on 498  degrees of freedom
## Residual deviance: 638.77  on 491  degrees of freedom
## AIC: 654.77
##
## Number of Fisher Scoring iterations: 4
```

```
#R squared and Wald test significance
require(rms)
lrn(logit.3)
```

```
## Logistic Regression Model
##
## lrm(formula = logit.3)
##
##              Model Likelihood      Discrimination      Rank Discrim.
##              Ratio Test              Indexes              Indexes
## Obs          499  LR chi2      20.31  R2          0.054  C          0.605
## 0            186  d.f.          7      g          0.481  Dxy         0.210
## 1            313  Pr(> chi2) 0.0049  gr          1.618  gamma        0.211
## max |deriv| 8e-06      gp          0.100  tau-a        0.098
##              Brier          0.226
##
##              Coef      S.E.   Wald Z Pr(>|Z|)
## Intercept      0.8597 0.9895   0.87 0.3849
## stabel        -0.2904 0.2656  -1.09 0.2742
## open          -0.0203 0.1034  -0.20 0.8446
## zorgvuldig    -0.4413 0.2519  -1.75 0.0797
## extravert      0.0766 0.1051   0.73 0.4662
## altruisme      0.0061 0.1217   0.05 0.9602
## age_opleiding  0.0556 0.0165   3.37 0.0008
## stabel * zorgvuldig 0.0738 0.0729   1.01 0.3111
##
```

```
#Interpret marginal effects
require(mfx)
logitmfx(logit.3, data = data.dropout)
```

```
## Call:
## logitmfx(formula = logit.3, data = data.dropout)
##
## Marginal Effects:
##              dF/dx  Std. Err.      z    P>|z|
## stabel          -0.0672418  0.0614541 -1.0942 0.273876
## open            -0.0046911  0.0239275 -0.1961 0.844566
## zorgvuldig      -0.1021752  0.0582265 -1.7548 0.079296 .
## extravert        0.0177345  0.0243313  0.7289 0.466077
## altrusme         0.0014058  0.0281838  0.0499 0.960217
## age_opleiding    0.0128629  0.0037718  3.4103 0.000649 ***
## stabel:zorgvuldig 0.0170968  0.0168691  1.0135 0.310823
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

*#Variance inflation factor: multicollinearity test*

```
require(car)
vif(logit.3)
```

```
##              stabel              open              zorgvuldig
##          10.899160           1.270075           8.289497
##          extravert          altrusme          age_opleiding
##          1.424272           1.184704           1.035412
## stabel:zorgvuldig
##          22.659327
```

```
sqrt(vif(logit.3)) > 2
```

```
##              stabel              open              zorgvuldig
##              TRUE              FALSE              TRUE
##          extravert          altrusme          age_opleiding
##              FALSE              FALSE              FALSE
## stabel:zorgvuldig
##              TRUE
```

*#Test for outliers*

```
outlierTest(logit.3)
```

```
##
## No Studentized residuals with Bonferonni p < 0.05
## Largest |rstudent|:
##          rstudent unadjusted p-value Bonferonni p
## 406 -1.949069           0.051287           NA
```

#### Fourth Logit Model

Dropped all the outliers for *age*: bigger than 35. Resulting in a sample size of 466 (before: 499).

```
logit.4 <- glm(vrv_1 ~ stabel + open + zorgvuldig + extravert + altrusme + age_opleiding,
               family = binomial(logit),
               data = subset(data.dropout, !age_opleiding > 35))
```

*#Display results*

```
summary(logit.4)
```

```
##
## Call:
```

```
## glm(formula = vrv_1 ~ stabiel + open + zorgvuldig + extravert +
##      altrusme + age_opleiding, family = binomial(logit), data = subset(data.dropout,
##      !age_opleiding > 35))
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5952  -1.3122   0.9136   1.0067   1.1694
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.457961   0.638476   0.717   0.4732
## stabiel       -0.045877   0.087652  -0.523   0.6007
## open          0.023058   0.103870   0.222   0.8243
## zorgvuldig    -0.171712   0.096692  -1.776   0.0758 .
## extravert      0.069962   0.106084   0.659   0.5096
## altrusme      -0.003223   0.123037  -0.026   0.9791
## age_opleiding  0.023985   0.026921   0.891   0.3730
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 624.39  on 465  degrees of freedom
## Residual deviance: 619.72  on 459  degrees of freedom
## AIC: 633.72
##
## Number of Fisher Scoring iterations: 4
```

```
#R squared and Wald test significance
require(rms)
lrm(logit.4)
```

```
## Logistic Regression Model
##
## lrm(formula = logit.4)
##
##              Model Likelihood    Discrimination    Rank Discrim.
##              Ratio Test          Indexes          Indexes
## Obs          466    LR chi2      4.67    R2          0.014    C          0.559
## 0             183    d.f.         6      g          0.236    Dxy         0.117
## 1             283    Pr(> chi2) 0.5869    gr          1.267    gamma        0.118
## max |deriv| 2e-08                gp          0.056    tau-a        0.056
##                               Brier          0.236
##
##              Coef    S.E.    Wald Z Pr(>|Z|)
## Intercept      0.4580 0.6385  0.72  0.4732
## stabiel       -0.0459 0.0877 -0.52  0.6007
## open          0.0231 0.1039  0.22  0.8243
## zorgvuldig    -0.1717 0.0967 -1.78  0.0758
## extravert      0.0700 0.1061  0.66  0.5096
## altrusme      -0.0032 0.1230 -0.03  0.9791
## age_opleiding  0.0240 0.0269  0.89  0.3730
##
```

```
#Marginal effects
require(mfx)
logitmfx(logit.4, data = data.dropout)

## Call:
## logitmfx(formula = logit.4, data = data.dropout)
##
## Marginal Effects:
##          dF/dx  Std. Err.      z    P>|z|
## stabel       -0.0084973  0.0201438 -0.4218 0.6731477
## open         -0.0024060  0.0238133 -0.1010 0.9195210
## zorgvuldig   -0.0479037  0.0222079 -2.1571 0.0310010 *
## extravert      0.0183749  0.0243173  0.7556 0.4498710
## altruisme      0.0026981  0.0281509  0.0958 0.9236436
## age_opleiding 0.0129064  0.0037613  3.4314 0.0006005 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#Variance inflation factor: multicollinearity test
require(car)
vif(logit.4)
```

```
##          stabel          open   zorgvuldig   extravert   altruisme
##          1.175143      1.249847      1.218655      1.411745      1.185207
## age_opleiding
##          1.037346
```

```
sqrt(vif(logit.4)) > 2
```

```
##          stabel          open   zorgvuldig   extravert   altruisme
##          FALSE          FALSE          FALSE          FALSE          FALSE
## age_opleiding
##          FALSE
```

```
#Test for outliers
outlierTest(logit.4)
```

```
##
## No Studentized residuals with Bonferonni p < 0.05
## Largest |rstudent|:
##          rstudent unadjusted p-value Bonferonni p
## 483 -1.617687      0.10573      NA
```

## Fifth Logit Model

Our definitive, kick-ass model.

```
logit.5 <- glm(vrv_1 ~ zorgvuldig + open + stabel + extravert + altruisme + zorgvuldig:allochtoonBin +
               family = binomial(logit),
               data = data.dropout)
```

```
#Display results
summary(logit.5)
```

```
##
## Call:
## glm(formula = vrv_1 ~ zorgvuldig + open + stabel + extravert +
```

```
##      altrusme + zorgvuldig:allochtoonBin + open:allochtoonBin +
##      stabiel:allochtoonBin + allochtoonBin + age_opleiding + leerintelligentie,
##      family = binomial(logit), data = data.dropout)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1111  -1.2039   0.7125   0.9641   1.6623
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -0.22638    0.68232  -0.332  0.740052
## zorgvuldig     -0.35826    0.11199  -3.199  0.001379 **
## open           0.09935    0.12336   0.805  0.420595
## stabiel        -0.15678    0.10259  -1.528  0.126463
## extravert       0.05850    0.10827   0.540  0.588996
## altrusme        0.05489    0.12679   0.433  0.665057
## allochtoonBinYes -1.56605    1.04383  -1.500  0.133540
## age_opleiding   0.05777    0.01711   3.377  0.000732 ***
## leerintelligentie 0.20454    0.11260   1.817  0.069283 .
## zorgvuldig:allochtoonBinYes 0.73863    0.25009   2.953  0.003143 **
## open:allochtoonBinYes -0.63185    0.24154  -2.616  0.008898 **
## stabiel:allochtoonBinYes 0.40362    0.21439   1.883  0.059748 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 659.08  on 498  degrees of freedom
## Residual deviance: 613.36  on 487  degrees of freedom
## AIC: 637.36
##
## Number of Fisher Scoring iterations: 4
```

```
#Calculate odds ratio vector
coef.vector.5 <- exp(logit.5$coef)

#R squared and Wald test significance
require(rms)
lrn(logit.5)
```

```
## Logistic Regression Model
##
##      lrn(formula = logit.5)
##
##              Model Likelihood      Discrimination      Rank Discrim.
##              Ratio Test      Indexes      Indexes
## Obs          499  LR chi2      45.72  R2          0.119  C          0.669
## 0            186  d.f.          11    g           0.774  Dxy         0.337
## 1            313  Pr(> chi2) <0.0001  gr          2.169  gamma        0.339
## max |deriv| 3e-05      gp          0.162  tau-a        0.158
##              Brier          0.214
##
##              Coef      S.E.  Wald Z Pr(>|Z|)
## Intercept    -0.2264 0.6823 -0.33  0.7401
## zorgvuldig    -0.3583 0.1120 -3.20  0.0014
```

```
## open 0.0994 0.1234 0.81 0.4206
## stabiel -0.1568 0.1026 -1.53 0.1265
## extravert 0.0585 0.1083 0.54 0.5890
## altrusme 0.0549 0.1268 0.43 0.6651
## allochtoonBin=Yes -1.5660 1.0438 -1.50 0.1335
## age_opleiding 0.0578 0.0171 3.38 0.0007
## leerintelligentie 0.2045 0.1126 1.82 0.0693
## zorgvuldig * allochtoonBin=Yes 0.7386 0.2501 2.95 0.0031
## open * allochtoonBin=Yes -0.6319 0.2415 -2.62 0.0089
## stabiel * allochtoonBin=Yes 0.4036 0.2144 1.88 0.0597
##
```

*#Marginal effects*

```
require(mfx)
```

```
mfx.5 <- logitmfx(logit.5, data = data.dropout)
```

*#Variance inflation factor: multicollinearity test*

```
require(car)
```

```
vif(logit.5)
```

```
##          zorgvuldig          open
##          1.517166          1.703091
##          stabiel          extravert
##          1.518127          1.454943
##          altrusme          allochtoonBinYes
##          1.214333          21.299938
##          age_opleiding          leerintelligentie
##          1.055990          1.260560
## zorgvuldig:allochtoonBinYes          open:allochtoonBinYes
##          18.119427          16.813727
##          stabiel:allochtoonBinYes
##          11.301315
```

```
sqrt(vif(logit.5)) > 2
```

```
##          zorgvuldig          open
##          FALSE          FALSE
##          stabiel          extravert
##          FALSE          FALSE
##          altrusme          allochtoonBinYes
##          FALSE          TRUE
##          age_opleiding          leerintelligentie
##          FALSE          FALSE
## zorgvuldig:allochtoonBinYes          open:allochtoonBinYes
##          TRUE          TRUE
##          stabiel:allochtoonBinYes
##          TRUE
```

*#Test for outliers*


```
outlierTest(logit.5)
```

```
##
## No Studentized residuals with Bonferonni p < 0.05
## Largest |rstudent|:
##          rstudent unadjusted p-value Bonferonni p
## 274 -2.167205          0.030219          NA
```

```
#Cook's distance
#plot(logit.5, which = 4, main = "Model Two")

ggplot(aes(x = seq_along(.cooks), y = .cooks), data = logit.5) +
  geom_bar(stat = "identity") +
  theme_bw(base_family = "serif") +
  labs(title = "Model Two: Cook's Distance",
       x = "Observation Number",
       y = "Cook's Distance") +
  geom_hline(yintercept = 0.0082, colour = 2, alpha = 0.75, linetype = 5)
```

Model Two: Cook's Distance



```
#Resid vs fitted
#plot(logit.5, which = 1, main = "Model Two")

ggplot(aes(x = .fitted, y = .resid), data = logit.5) +
  geom_point(alpha = 1/4) +
  geom_hline(yintercept = 0, linetype = 2, alpha = 0.75) +
  geom_smooth(alpha = 0.15, method = "loess") +
  theme_bw(base_family = "serif") +
  labs(title = "Model Two: Residuals vs. Fitted",
       x = "Predicted Values",
       y = "Residuals")
```

Model Two: Residuals vs. Fitted

